

## Case 2: Creating Data as a Service

### Instructions:

- You can use either Python or R to work on this case.
- No sharing of work. You can work in your teams only
- You are expected to submit a report that summarizes the key steps in your implementation as a flow chart and submit fully functional code.
- Deadline: 07/07/2017 11.59 PM. Late submissions lose 10% points per day.

### Working with large datasets:

We are going to expand the housing problem we worked in the class and work with real Zillow data.

**Review** <https://www.continuum.io/blog/developer-blog/productionizing-and-deploying-data-science-projects> for motivation on why we are working on this assignment

### Review the following to understand the problem:

- <https://www.kaggle.com/c/zillow-prize-1>
- <https://www.kaggle.com/philippsp/exploratory-analysis-zillow>
- <https://www.kaggle.com/sudalairajkumar/simple-exploration-notebook-zillow-prize>
- <https://www.kaggle.com/captcalculator/a-very-extensive-zillow-exploratory-analysis>

### 1. Data ingestion, EDA, Wrangling:

- Download the data from Zillow. (<https://www.kaggle.com/c/zillow-prize-1>)
- Create an IPYB notebook and Conduct an in-depth EDA (See below for ideas; Note: Your code should be original. You are welcome to use ideas but with attribution).
- Put together a note on what data cleansing is required for automation
- Clean up the data and take care of missing data values using a Python/R script
- Programmatically write the data to a S3 bucket named "ZillowData". (Use a configuration file to put your Amazon keys)

```
{ "AWSAccess": "access key",
  "AWSSecret": "secret key"}
```

- Dockerize this image: Your folder structure should be something like this (assuming all data you are uploading is in the data folder):
- Assignment2
  - config.json
  - run.sh
  - dataIngestion.py
  - rawDataEDA.ipynb
  - data
  - ddmmyyMMSS.log
- Submit your image to dockerHub and upload all files to github.

## 2. Create a DBaas (Database as a service)

- Using the DBaas assigned to you, move the clean data to a cloud-based database.
- Note: You should use command line to do this. Create a script that will take the clean data you created earlier and create a DBaas. You will have to review the apis for your assigned DBaas to do this.
- Write a Jupyter notebook illustrating how you can run sample queries connecting to the cloud DBaas and getting back the data.

### Data storage:

1. For this exercise, you will be saving the clean data in the system assigned to you
  - a. Team 1: Amazon RDS
  - b. Team 2 : Amazon S3 – Use Athena
  - c. Team 3 : Amazon SimpleDB
  - d. Team 4 : Google BigQuery
  - e. Team 5 : IBM Cloudant
  - f. Team 6 : Microsoft SQL Server db on the cloud
  - g. Team 7 : MongoDB Atlas
  - h. Team 8 : MongoDB on an EC2 instance
  - i. Team 9 : Postgres SQL on an EC2 instance
  - j. Team 10 : SQLite on an EC2 instance

## 3. Create a Rest API to serve the data:

- Your next task is to create a Web services REST API to serve data
- Preferably use FLASK to do this.
- Your web service must be hosted on the cloud (Choose any)
- Here are a few example articles you should review to get a good idea on how you can go about doing this:
  - <https://medium.com/@joeclark.phd/creating-a-data-product-with-flask-mongodb-and-ibm-bluemix-7deada61c573>
  - <http://blog.luisrei.com/articles/flaskrest.html>
  - <https://blog.threatstack.com/writing-a-web-service-using-python-flask>
  - <https://blog.miguelgrinberg.com/post/designing-a-restful-api-with-python-and-flask>
  - <http://blog.thedataincubator.com/2015/09/painlessly-deploying-data-apps-with-bokeh-flask-and-heroku/>
  - <https://www.codementor.io/jadianes/building-a-web-service-with-apache-spark-flask-example-app-part2-du1083854>
  - <https://content.pivotal.io/blog/data-science-how-to-text-analytics-as-a-service>
- Create a Jupyter notebook to illustrate how to use your REST API

#### 4. Enhancing your REST API: Geospatial search

Note that each record has a Latitude and Longitude. Your goal is to create a REST API that given a Lat and Long, should return the 10 closest homes.

Write a Jupyter notebook and illustrate using this REST API. You should present results something like this

```
+-----+-----+-----+-----+
| hotel_name | lat | lon | dist |
+-----+-----+-----+-----+
| Hotel Astori.. | 122.41 | 37.79 | 0.0054 |
| Juliana Hote.. | 122.41 | 37.79 | 0.0069 |
| Orchard Gard.. | 122.41 | 37.79 | 0.0345 |
| Orchard Gard.. | 122.41 | 37.79 | 0.0345 |
...
+-----+-----+-----+-----+
10 rows in set (4.10 sec)
```

..

Review these articles for the algorithm and how to use SQL to get these results:

- [http://www.arubin.org/files/geo\\_search.pdf](http://www.arubin.org/files/geo_search.pdf)
- <https://www.percona.com/blog/2014/06/19/using-udfs-for-geo-distance-search-in-mysql/>
- <https://www.percona.com/blog/2013/10/21/using-the-new-mysql-spatial-functions-5-6-for-geo-enabled-applications/>